

Critique of the AI Collaborator

Throughout this assignment the AI was a powerful accelerator, but it was not infallible, and several concrete failures taught me where its judgment breaks down. Its most revealing mistake was **wrong**: a probe titled “admin cancels a shipping order” actually executed a *user* cancellation, so the `200` it reported was really a different, already-known defect (BUG-07). The AI produced this case by pattern-matching plausible-looking code rather than verifying that the request matched its own description, and the error only surfaced when I questioned the result. A related **bias** was over-compliance: when I casually said an ambiguous case “is a bug,” it immediately filed a GitHub issue (BUG-14) claiming admin *should* be able to cancel a shipping order, even though the state diagram made the existing behaviour correct — we had to revert the issue. It also defaulted several `spec-undefined` probes and the coupon case-sensitivity finding to “not a bug,” reflecting a conservative lean toward whatever the specification literally said rather than reasoning about consistency. Its work was also **incomplete**: its own FR-10 mobile redesign silently dropped admin-transition coverage until I explicitly asked whether more bugs remained, and the `.http` file it wrote used a token-variable indirection that broke admin login because the AI never actually ran the tool it generated. The common thread is that the AI is fast and confident but does not reliably check its own artifacts or resist ambiguity — it optimises for producing a plausible answer, not for being certain. The principle I learned is that effective AI collaboration requires the human to remain the **oracle and the skeptic**: I must verify the AI’s generated tests, tooling, and classifications against the running system, treat its confidence as a hypothesis rather than a verdict, and never let its eagerness to comply replace my own judgment on ambiguous requirements.